

# CPSC 457

## Assignment #1

Ahmad Almasri - 30114233

# Question 1

t3.txt

```
ahmad.almasri3@gfx03-7:~/cp457/a1$ time python3 palindrome.py < t3.txt
Longest palindrome: __o.O.o__

real    0m0.025s
user    0m0.015s
sys     0m0.005s
```

t4.txt

```
ahmad.almasri3@gfx03-7:~/cp457/a1$ time python3 palindrome.py < t4.txt
Longest palindrome: redder

real    0m0.260s
user    0m0.241s
sys     0m0.011s
```

t3.txt

t4.txt

```
ahmad.almasri3@gfx03-7:~/cpsc457/a1$ time ./slow-pali < t3.txt
Longest palindrome: __o.o.o__

real    0m0.009s
user    0m0.002s
sys     0m0.003s
```

```
ahmad.almasri3@gfx03-7:~/cpsc457/a1$ time ./slow-pali < t4.txt
Longest palindrome: redder

real    0m2.454s
user    0m1.116s
sys     0m1.334s
```

## Part b

---

t3.txt		
	<b>C++</b>	<b>Pyhon</b>
User	0.002s	0.015s
Kernel	0.003s	0.005s

t4.txt		
	<b>C++</b>	<b>Pyhon</b>
User	1.116s	0.241s
Kernel	1.334s	0.011s

In the first place, there is a main difference between C++ and Python. C++ is considered to be a **compiled language**. However, Python is a **compiled and interpreted** one.

**Compiled language** means that we compile the source code to a **machine language** (binary file). The binary file is called the **executable file**. Because of the binary file, running the c++ code would be very **fast** because the code would be compiled ahead of time. **However**, having a compiled file requires us to have different versions for each OS.

Python is an **intermediate language** (Bytecode). Where, the source code of python being compiled to a **Bytecode**. The Bytecode would be interpreted line by line, while executing the program. For example, the interpreter for Python is **PVM** and JAVA is the **JVM**.

**In case of t3.txt**, C++ was faster than Python. I think, the reason of that result is that the size of the text file is small. On the other hand, **the t4.txt** was bigger in size. Python reads the whole content by using less number of system calls.

In general, C++ is faster than Python, but the way that we write the code must be efficient to take advantage of the language either C++ or Python.

## Question 2

```
ahmad.almasri3@gfx-ta3:~/cp457/a1$ ./dup.py 200000000 < t4.txt | time ./fast-pali
Longest palindrome: redder
24.34user 0.27system 0:25.53elapsed 96%CPU (0avgtext+0avgdata 3980maxresident)k
0inputs+0outputs (0major+383minor)pagefaults 0swaps
ahmad.almasri3@gfx-ta3:~/cp457/a1$
```



## Question 3

```
ahmad.almasri3@gfx-ta3:~/cp457/a1$ strace -c ./fast-pali < t4.txt
```

```
Longest palindrome: redder
```

% time	seconds	usecs/call	calls	errors	syscall
50.37	0.001309	72	18		read
21.43	0.000557	11	48	43	openat
12.31	0.000320	14	22		mmap
3.66	0.000095	13	7		mprotect
3.27	0.000085	10	8	7	stat
2.27	0.000059	8	7		lseek
1.77	0.000046	7	6		fstat
1.69	0.000044	8	5		close
1.08	0.000028	9	3		brk
0.85	0.000022	22	1		munmap
0.65	0.000017	17	1	1	access
0.65	0.000017	8	2	1	arch_prctl
0.00	0.000000	0	1		write
0.00	0.000000	0	1		execve
100.00	0.002599	19	130	52	total

```
ahmad.almasri3@gfx-ta3:~/cp457/a1$ strace -c ./fast-pali < t3.txt
```

```
Longest palindrome: __o.O.o__
```

% time	seconds	usecs/call	calls	errors	syscall
28.41	0.000125	2	48	43	openat
24.32	0.000107	107	1		execve
17.05	0.000075	3	22		mmap
6.59	0.000029	2	13		read
5.68	0.000025	3	7		mprotect
4.09	0.000018	2	8	7	stat
2.95	0.000013	1	7		lseek
2.73	0.000012	2	6		fstat
2.27	0.000010	2	5		close
1.59	0.000007	7	1		munmap
1.59	0.000007	2	3		brk
0.91	0.000004	4	1		write
0.91	0.000004	4	1	1	access
0.91	0.000004	2	2	1	arch_prctl
100.00	0.000440	3	125	52	total

```
ahmad.almasri3@gfx-ta3:~/cpssc457/a1$ strace -c ./slow-pali < t4.txt
```

```
Longest palindrome: redder
```

% time	seconds	usecs/call	calls	errors	syscall
100.00	9.592183	1	5767205		read
0.00	0.000075	1	48	43	openat
0.00	0.000053	7	7		mprotect
0.00	0.000052	2	22		mmap
0.00	0.000009	1	7		lseek
0.00	0.000008	8	1		write
0.00	0.000008	1	6		fstat
0.00	0.000006	1	5		close
0.00	0.000004	4	1		munmap
0.00	0.000002	0	3		brk
0.00	0.000002	1	2	1	arch_prctl
0.00	0.000000	0	8	7	stat
0.00	0.000000	0	1	1	access
0.00	0.000000	0	1		execve
100.00	9.592402	1	5767317	52	total

## Part a

```
ahmad.almasri3@gfx-ta3:~/cp457/a1$ strace -c ./slow-pali < t3.txt
```

```
Longest palindrome: __o.O.o__
```

% time	seconds	usecs/call	calls	errors	syscall
32.74	0.000204	204	1		execve
20.55	0.000128	2	48	43	openat
16.69	0.000104	2	50		read
12.20	0.000076	3	22		mmap
4.01	0.000025	3	7		mprotect
3.21	0.000020	2	8	7	stat
2.09	0.000013	2	6		fstat
2.09	0.000013	1	7		lseek
1.61	0.000010	2	5		close
1.28	0.000008	2	3		brk
1.12	0.000007	7	1		munmap
0.96	0.000006	6	1		write
0.80	0.000005	5	1	1	access
0.64	0.000004	2	2	1	arch_prctl
100.00	0.000623	3	162	52	total

fast-pali.cpp is way faster than slow-pali.cpp. The main reason is the reduced number of SYS calls that we are performing to read the data. The size of the buffer in the slow version is only 1 Byte; however, the fast version is 1 MB.

## Part b

```
ahmad.almasri3@gfx-ta3:~/cpsc457/a1$ strace -c python3 palindrome.py < t4.txt
Longest palindrome: redder
% time      seconds  usecs/call   calls   errors syscall
-----
 17.61    0.000200         1    175     47 stat
 13.91    0.000158         1    141     76 openat
 12.24    0.000139         0    788      read
  9.77    0.000111         1    100      fstat
  8.98    0.000102         1     58      mmap
  8.80    0.000100         9     11      mprotect
  7.66    0.000087         1     68      close
  7.22    0.000082         5     16      getdents64
  3.26    0.000037         0     42      lseek
  2.46    0.000028         1     18     11 ioctl
  2.38    0.000027         0     58      brk
  1.14    0.000013        13      1      lstat
  0.88    0.000010         3      3      dup
  0.79    0.000009         4      2      munmap
  0.53    0.000006         2      3      2 readlink
  0.53    0.000006         3      2      futex
  0.35    0.000004         0     68      rt_sigaction
  0.26    0.000003         1      3      fcntl
  0.26    0.000003         3      1      getrandom
  0.18    0.000002         2      1      rt_sigprocmask
  0.18    0.000002         1      2      1 arch_prctl
  0.18    0.000002         2      1      set_tid_address
  0.18    0.000002         2      1      set_robust_list
  0.18    0.000002         2      1      prlimit64
  0.09    0.000001         1      1      getcwd
  0.00    0.000000         0      1      write
  0.00    0.000000         0      1      1 access
  0.00    0.000000         0      1      getpid
  0.00    0.000000         0      1      execve
  0.00    0.000000         0      1      sysinfo
  0.00    0.000000         0      1      getuid
  0.00    0.000000         0      1      getgid
  0.00    0.000000         0      1      geteuid
  0.00    0.000000         0      1      getegid
  0.00    0.000000         0      3      sigaltstack
-----
100.00    0.001136         0   1577    140 total
```

## Part b

```

ahmad.almasri3@gfx-ta3:~/cp457/a1$ strace -c python3 palindrome.py < t3.txt
Longest palindrome: __o.o.o__
% time      seconds  usecs/call   calls   errors syscall
-----
19.79      0.000285         1    175     47 stat
19.65      0.000283         2    141     76 openat
14.93      0.000215        13     16      getdents64
10.21      0.000147         2     58      mmap
10.07      0.000145         1     84      read
 7.08      0.000102         1    100      fstat
 6.46      0.000093         1     68      close
 2.85      0.000041         0     42      lseek
 2.01      0.000029         2     11      mprotect
 1.74      0.000025        25      1      lstat
 1.18      0.000017         0     18     11 ioctl
 0.83      0.000012         0     18      brk
 0.62      0.000009         3      3      2 readlink
 0.35      0.000005         5      1      write
 0.28      0.000004         0     68      rt_sigaction
 0.28      0.000004         1      3      fcntl
 0.28      0.000004         1      3      sigaltstack
 0.28      0.000004         2      2      1 arch_prctl
 0.21      0.000003         3      1      1 access
 0.21      0.000003         1      3      dup
 0.21      0.000003         3      1      getcwd
 0.21      0.000003         1      2      futex
 0.14      0.000002         2      1      getrandom
 0.07      0.000001         1      1      getuid
 0.07      0.000001         1      1      getegid
 0.00      0.000000         0      2      munmap
 0.00      0.000000         0      1      rt_sigprocmask
 0.00      0.000000         0      1      getpid
 0.00      0.000000         0      1      execve
 0.00      0.000000         0      1      sysinfo
 0.00      0.000000         0      1      getgid
 0.00      0.000000         0      1      geteuid
 0.00      0.000000         0      1      set_tid_address
 0.00      0.000000         0      1      set_robust_list
 0.00      0.000000         0      1      prlimit64
-----
100.00     0.001440         1    833    140 total

```



Comparing it with the python code, the c++ version is faster. First, c++ is faster than python in general. Second, the size of the buffer significantly reduced the SYS calls for fast-pali.cpp.